

## I. High-Level Architectural Flow

The system is built on an **Event-Driven Architecture (EDA)**. This means the frontend (Dashboard) and backend (Processing) are decoupled. The "glue" holding them together is the Kafka message bus (Confluent Cloud).

---

## II. Detailed Technical Flow

### Phase 1: Identity & Entitlements (Steps 1–3)

**Goal:** Establish *who* the user is and *what* they are allowed to do without hardcoding logic in the application code.

- **1. Authentication (Clerk / Auth0):**
  - User signs up on the frontend.
  - **Result:** User gets a JWT (JSON Web Token) containing sub: "user\_123". This token is sent with every subsequent API request.<sup>1</sup>
- **2. Entitlement Service (The Gatekeeper):**
  - **Database:** A table in **Neon (Postgres)** called user\_subscriptions.
  - **Workflow:**
    1. User buys a plan via Stripe on the frontend.
    2. Stripe sends a webhook (checkout.session.completed) to your Go backend.
    3. Backend updates Neon:  
**SQL**  
`INSERT INTO user_subscriptions (user_id, plan, features, expiry)  
VALUES ('user_123', 'ADVANCED', '{"realtime": true, "alerting": true}', '2026-01-01');`
  - **Frontend Check:** When the dashboard loads, it calls GET /v1/user/me. The backend joins the Clerk ID with the Neon subscription table. If plan == 'FREE', the "API Keys" tab is grayed out.

### Phase 2: The Unified Ingestion Layer (Steps 4 & 8)

**Goal:** Treat Batch (CSV) and Streaming (API) as the *same* data source for the detection engine.

- **Path A: Batch Ingestion (Async CSV):**
  1. **Upload:** User POSTs a file to the backend.
  2. **Storage:** Backend streams the file to **Cloudflare R2** (S3 compatible) to avoid memory crashes.
  3. **Hydration:** A background Goroutine reads the file from R2 line-by-line.
  4. **Publish:** It validates each row and pushes it to the **Confluent Kafka** topic transactions.
  5. **User Feedback:** The API returns 202 Accepted immediately. A WebSocket or polling

mechanism updates the UI progress bar.

- **Path B: Real-Time Ingestion (API - Advanced Plan):**
  1. **Ingress:** Go-based service receives POST /transaction.
  2. **Entitlement Check:** Checks **Upstash Redis** cache: GET subscription:user\_123. If invalid, return 403 Forbidden.
  3. **Publish:** Pushes directly to the *same* transactions topic used by the CSV worker.

**Technical Win:** The downstream ML engine doesn't know if a transaction came from a CSV or a live API call. It processes everything uniformly.

### Phase 3: The Detection Engine (Step 5)

**Goal:** Apply hybrid detection logic (Rules + ML) in parallel.

- **The Message Bus (Kafka Topic: transactions):** All data lands here first.
- **Step 5.1 & 5.2: Enrichment & Rules (Stream Processor):**
  - **Technology:** A Python worker (using confluent-kafka and pandas for windowing) or a lightweight Flink job.
  - **Stateful Processing:** It maintains a rolling window of the last 10 minutes of transactions for this user.
  - **Feature Lookup:** It queries **Upstash Redis** for historical baselines (e.g., user\_avg\_spend).
  - **Rule Execution:** It runs deterministic checks (Velocity, Geo-fencing).
  - **Output:** Adds metadata to the event: {"rules\_triggered": ["VELOCITY\_HIGH"]}.
- **Step 5.3: ML Inference (The Model):**
  - **Service:** A Python consumer reads the enriched stream.
  - **Model:** Uses **ONNX Runtime** to run a pre-trained Isolation Forest model. This is lightweight and runs inside the container.
  - **Output:** Generates a probability score (e.g., 0.92).
  - **Non-Blocking:** This happens asynchronously; it does not block the ingestion API.
- **Step 5.4 & 5.5: The Decision & Explainability (Aggregation):**
  - **Logic:** A final "Aggregator" function combines the Rule Output and ML Score.
  - **Explainability Generator:** It maps error codes to human strings.
    - *Input:* rule: VELOCITY\_5\_MIN, ml\_feature\_contribution: [amount, time]
    - *Output:* "High velocity detected: 5th transaction in 5 minutes. Amount (\$5000) is significantly above user average (\$150)."

### Phase 4: Storage & Presentation (Steps 6–7)

**Goal:** Efficiently store data for the dashboard.

- **Database (Neon Postgres):**
  - We use a standard Postgres table, partitioned by date if possible.
  - Schema: transactions table stores the raw data, anomalies table stores the detected issues with the JSON explanation.
- **Dashboard API:**

- The frontend calls GET /v1/anomalies.
- The backend queries Neon: SELECT \* FROM anomalies WHERE user\_id = 'user\_123' ORDER BY timestamp DESC.
- Since Neon is serverless, it scales down when you aren't developing, costing \$0.

## Phase 5: Action & Evolution (Steps 9–10)

**Goal:** Proactive alerts and self-improving models.

- **Alerting Engine (The Notification Service):**
  - **Consumer:** Listens to the anomalies Kafka topic.
  - **Filter:** Checks user config (if severity == HIGH and plan == ADVANCED).
  - **Deduping:** Uses Redis keys (e.g., SETNX alert:user\_123:last\_10\_min) to prevent spamming the user.
  - **Delivery:** Calls a free email provider API (like SendGrid free tier) or logs to console for dev.
- **Model Retraining (The Feedback Loop):**
  - **Trigger:** A weekly script.
  - **Data Source:** Pulls last week's data from Neon.
  - **Process:** Retrains the Isolation Forest model locally or in a cloud function.
  - **Deploy:** Saves the new .onnx model file to R2 storage. The running services pull the new model on restart.

## III. System Schema (Core Data Structure)

This is the JSON object that flows through your Kafka pipeline.

JSON

```
{
  "meta": {
    "trace_id": "abc-123-xyz",
    "timestamp": "2025-12-20T12:00:00Z",
    "source": "REALTIME_API", // or "BATCH_CSV"
    "user_id": "user_99"
  },
  "data": {
    "tx_id": "tx_555",
    "amount": 4500.50,
    "currency": "USD",
    "location": "NY"
  }
}
```

```
 },
  "enrichment": {
    "user_avg_spend": 120.00,
    "distance_from_last_tx": 5000 // km
  },
  "analysis": {
    "rule_flags": ["GEO_IMPOSSIBLE", "AMOUNT_SPIKE"],
    "ml_score": 0.98,
    "ml_prediction": "ANOMALY"
  },
  "verdict": {
    "final_severity": "CRITICAL",
    "explanation": "Transaction amount is 37x higher than average. Location (NY) is physically impossible relative to previous transaction (London, 10 mins ago)."
  }
}
```

## IV. Summary of the Free Cloud Stack

- **Identity:** Clerk (Free tier)
- **Ingestion/Queue:** Confluent Cloud (Kafka) (Free credit)
- **Cache/State:** Upstash Redis (Free tier)
- **Storage/DB:** Neon Postgres (Free tier) & Cloudflare R2 (Free tier)
- **Compute:** Your Laptop (Development) -> Render/Railway (Deployment - Free/Cheap tiers)